

# Namespaces

# Namespaces

- namespaces of the current context

`ls -l /proc/self/ns` or `ls -l /proc/$$/ns`

- cgroup control / limitation of resources
- ipc hmessage queues, semaphores, shared mem for ipc
- mnt mount
- net network
- pid, pid\_for\_children process id
- user user (UIDs, GIDs)
- uts hostname & domainname
- time time

# Network Namespaces

- provide the isolation of networking related resources (see `man network_namespaces.7`)
- network devices,
- IPv4 and IPv6 protocol stacks,
- IP routing tables, firewall rules,
- port numbers (sockets)
- some filesystem entities
  - `/proc/net` directory (a symbolic link to `/proc/PID/net`),
  - `/sys/class/net` directory,
  - various files under `/proc/sys/net`,

# Network Namespaces & Device Configuration

- some linux commands in network namespace context
  - `ip netns list`
  - `ip netns add <netns_name> or`  
`ip netns delete <netns_name>`
  - `ip netns exec <netns_name> <command>`
  - `ip link show (or ifconfig or ifconfig -a)`
  - `ip link add <veth_name1> type veth peer name`  
`<veth_name2>`
  - `ip link set <veth_name> netns <netns_name>`
  - `ip link set dev <veth_name> up (or down)`
  - `ip addr add <ip_addr>/<net_bits> dev <eth_name>`

# Network Namespaces & Device Configuration

- some linux commands in bridging context
  - `brctl show`
  - `brctl addbr <br_name>`
  - `brctl stp <br_name> on`
  - `brctl setageing <br_name> <time> (0 ↔ off)`
  - `ip link set dev <br_name> up`
  - `ip neigh {show|add|del| ...}`
  - `(old: arp -nv , arp -nd <ip_addr>)`
  - `vconfig add <vlan_name> <vlan_no>`
  - `ip link set dev <vlan_name.vlan_no> up`  
(or down)

# Namespaces

- namespaces of the current context

`ls -l /proc/self/ns` or `ls -l /proc/$$/ns`

- cgroup    control / limitation of resources
- ipc    hmessage queues, semaphores, shared mem for ipc
- mnt    mount
- net    network
- pid, pid\_for\_children process id
- user    user (UIDs, GIDs)
- uts    hostname & domainname
- time    time

# Mount Namespaces

- some linux commands in mount namespace context
  - `findmnt` display mountpoints
  - `sudo find /proc/*/ns -name 'mnt' -exec readlink {} \; | sort -u` find all mnt namespaces
  - `unshare -m [<prog>]` create a new mnt namespace and run <prog> within (or run `$SHELL` by default)
  - `mount --bind d1 d2` make the same content of d1 available at the mountpoint d2.  
Works for files too (file d2 already has to exist).  
Recursive operation for submounts by `--rbind`.  
Make a directory a mountpoint by `mount --bind dir dir`  
Property change by remounting, e.g. make d2 read only while d1 is still writeable by  
`mount -o remount,bind,ro d1 d2`
  - `lsns` display namespaces and processes

# Mount Namespaces

- playing around: creating / deleting
  - list all mnt namespaces
  - create a new one by issuing `(sudo) unshare -m`
  - again list all mnt namespaces and observe the newly created
  - exit the shell
  - again list all mnt namespaces and observe that the newly created one has vanished



# Mount Namespaces

- playing around: isolated fs
  - create a shell in a new mnt namespace `sudo unshare -m`
  - observe existing mountpoints from within the new ns and from outside using `findmnt` → the same
  - create a new temporary filesystem within the namespace  
`mkdir r1`  
`mount -t tmpfs tmpfs r1`
  - again observe existing mountpoints from within the new ns and from outside using `findmnt` → the `r1` mountpoint is only available in the the new ns.
  - create a file in the shared filesystem and observe its ownership from within the new ns and from outside → both root
  - repeat the experiment creating a new mnt namespace by `unshare -Urm` and observe the file ownership → user / root

# Mount Namespaces

- playing around: isolated root fs
  - get the miniroot filesystem from  
<https://alpinelinux.org/downloads/>
  - create the location for the mnt ns new root fs and prepare it  
`mkdir my_root`  
`tar -xf alpine...tar.gz -C my_root`
  - create the new mnt ns and make the future root fs a mountpoint (by the way create a user ns (-U) and map root within the ns to vagrant outside (-r))  
`unshare -Urm (now we are in our new mnt ns)`  
`mount --bind my_root my_root`

# Mount Namespaces

- ... playing around: isolated root fs
  - make a directory `old_root` to map the current root fs to  
`cd my_root && mkdir old_root`
  - change the shell to `sh` (alpine does not provide `bash`, so after the root fs change `bash` is no longer available)  
`sh`
  - exchange the old and the new root fs .. that's it!  
`pivot_root . old_root`  
`cd /`
  - clean up (dissolve the connection to the outer fs)  
`umount -l old_root`
  - now in the `mnt ns` we have an miniature system without access other parts of the VMs filesystem and with its own executables (based on busybox). From the outer VM we can look into `my_root`, where `mnt ns`' / is living in.

# Mount Namespaces

- playing around: isolated overlay fs
  - create the new mnt ns  
`unshare -m`
  - create a directory `over_dir` to hold the contents to lay over `lowerdir`, and a working directory `over_work`  
`mkdir over_dir && mkdir over_work`
  - create the overlay  
`mount -t overlay -o lowerdir=/etc, \`  
`upperdir=./over_dir,workdir=./over_work \`  
`overlayfs /etc`
  - from within the mnt ns writes to `/etc` effectively go to `over_dir`, and vice versa
  - thanks to the mnt ns changes within the mnt ns have no effect on the outside (`/etc` remains unaltered). Outside changes to the (underlying) `/etc` are visible within the mnt ns too.

# Namespaces

- some (more) useful commands
  - (sudo) `lsns -t net`  
display the (network) namespaces and (only) the 1st process under each namespace
  - (sudo) `ps -e -o netns,pid,cmd --sort netns`  
display processes under a (network) namespace
  - (sudo) `nsenter -t <pid> -a <cmd>`  
execute <cmd> (default is \$SHELL) under all (-a) the same namespaces as the target process with <pid>
  - (sudo) `ip link add veth1 netns <name1|pid1> \`  
`type veth peer name veth2 netns <name2|pid2>`  
add a virtual ethernet connection's endpoints directly into the namespaces identified by their names <name1> and <name2> or the processes' <pid1> and <pid2> that belong to them.